



Microaneurysm detection using fully convolutional neural networks

Piotr Chudzik^{a,*}, Somshubra Majumdar^b, Francesco Calivá^a, Bashir Al-Diri^a, Andrew Hunter^a

^a School of Computer Science, University of Lincoln, Lincoln LN6 7TS, UK

^b Department of Computer Science, University of Illinois, Chicago, IL 60607, USA

ARTICLE INFO

Article history:

Received 3 July 2017

Revised 18 January 2018

Accepted 22 February 2018

Keywords:

Medical image analysis

Microaneurysm detection

Convolutional neural networks

Retinal fundus images

ABSTRACT

Background and Objectives: Diabetic retinopathy is a microvascular complication of diabetes that can lead to sight loss if treated not early enough. Microaneurysms are the earliest clinical signs of diabetic retinopathy. This paper presents an automatic method for detecting microaneurysms in fundus photographs.

Methods: A novel patch-based fully convolutional neural network with batch normalization layers and Dice loss function is proposed. Compared to other methods that require up to five processing stages, it requires only three. Furthermore, to the best of the authors' knowledge, this is the first paper that shows how to successfully transfer knowledge between datasets in the microaneurysm detection domain.

Results: The proposed method was evaluated using three publicly available and widely used datasets: E-Ophtha, DIARETDB1, and ROC. It achieved better results than state-of-the-art methods using the FROC metric. The proposed algorithm accomplished highest sensitivities for low false positive rates, which is particularly important for screening purposes.

Conclusions: Performance, simplicity, and robustness of the proposed method demonstrates its suitability for diabetic retinopathy screening applications.

© 2018 Elsevier B.V. All rights reserved.

1. Introduction

Diabetes affects one in eleven adults (over 400 million people worldwide) [1]. Diabetic retinopathy (DR) is a microvascular complication of diabetes which is the leading cause of vision loss in the working-age population [2]. One out of three diabetics has DR [3] and one in ten diabetic patients develops most vision-threatening form of DR [4]. Early detection of DR can prevent blindness in 90% of cases [5].

DR screening is manually performed by ophthalmologists and trained graders through a visual inspection of fundus photographs (FP). Unfortunately, the grading process is time-consuming, tedious, and error-prone with high inter-observer variability. Due to the rising number of DR patients worldwide (expected to exceed 640m by 2040 [1]) and their location (75% live in underdeveloped areas [6]) the development of computer-assisted diagnosis and automatic DR screening approaches are of the utmost importance.

Microaneurysms (MAs) are spherical swellings of the capillaries caused by weakening of the vascular walls; they appear as small round red dots. They are the earliest clinical sign of DR and continue to be present as the disease progresses. Consequently,

automated detection of MAs can drastically reduce the screening workload. MA detection is a challenging task even for the human eye due to many factors including uneven image illumination, reflections, limited resolution and media opacity. The boundaries of MAs are not always well-defined and local contrast to the background is low, even in high-resolution images. Moreover, MAs may be confounded with visually similar anatomical structures such as haemorrhages, junctions in thin vessels, disconnected vessel segments, dark patches on vessels, background pigmentation patches and dust particles on the camera lense.

In general, the majority of MA detection methods consists of up to five stages: 1) Preprocessing, 2) MA candidate extraction, 3) Vessel removal, 4) Candidate feature extraction, and 5) Classification. The main goal of preprocessing is to remove noise, correct non-uniform illumination, and to improve contrast between the MAs and background. The MA candidate extraction stage uses a simple algorithm to identify a reasonably small set of locations with somewhat “lesion-like” appearance, attempting to identify all actual lesions together with many false positive regions. The vessel removal stage addresses the large number of false positives that may otherwise be produced by vessels. Next, hand-crafted features are extracted from candidate regions; this is the most labour-intensive and time-consuming part of the design stage. Finally, a classifier is trained to distinguish MAs from non-MAs based on the extracted features.

* Corresponding author.

E-mail address: plukasz@lincoln.ac.uk (P. Chudzik).

Baudoin et al. [7] introduced the first MA detection algorithm applied to fluorescein angiogram images. They employed a mathematical morphology based approach to remove vessels and applied a top-hat transformation with linear structuring elements to detect MAs. Several methods were built on this approach [8], however, since intravenous use of fluorescein can cause death in 1 in 222 000 cases [9], such methods are not suited for screening purposes. Walter et al. [10] also used a top-hat based method and automated thresholding to extract MA candidates. They extracted 15 features and applied kernel density estimation with variable bandwidth for MA classification. In general, morphology-based approaches are sensitive to changes in size and shape of structuring elements which result in significant variations in MAs detection results. Zhang et al. [11] proposed a method based on dynamic thresholding and correlation coefficients of a multi-scale Gaussian template. They used 31 manually designed features based on intensity, shape and response of a Gaussian filter. Veiga et al. [12] presented an algorithm using Law texture features. Support Vector Machines (SVM) were used in a cascading manner: first SVM was used to extract MA candidates whereas the second SVM performed final MA classification. Haloi [13] used a vanilla convolutional network with 3 convolutional layers and 2 fully connected layers to detect MAs. Javidi et al. [8] proposed a technique which used 2D Morlet wavelet to find MA candidates. At the next stage, a discriminative dictionary learning approach was employed to distinguish MAs from other structures. Srivastava et al. [14] used Frangi-based filters that were manually designed to distinguish vessels from red lesions. Filters were applied to multiple sized image patches to extract features. Finally, these features were classified using a SVM.

Compared to the methods mentioned above, the proposed algorithm requires only three stages instead of five (preprocessing, patch extraction and classification). There is no need for MA candidate detection, vessel removal or feature extraction. Furthermore, the proposed method does not require manually hand-crafted features, it automatically learns the most discriminative features for MA detection. The vast majority of MA detection algorithms employ features based on MA shape, colour and texture. Unfortunately, many image modalities makes it virtually impossible to model them manually. To address this challenge, a Convolutional Neural Network (CNN) was used. CNNs have emerged as a powerful family of algorithms for solving computer vision tasks such as object detection [15], semantic segmentation [16] and image classification [17]. Compared with Haloi [13] method, the presented algorithm proposes a novel fully convolutional neural network (FCNN) architecture and transfers knowledge between MA datasets.

Training CNNs from scratch is not a trivial task, as they require large amounts of labelled data for training. In the MA detection domain, public datasets are small, scarce, and local lesion annotations on a per-pixel level are almost non-existent (to the best of authors knowledge, only one such dataset exists [18]). Moreover, the CNNs have vast capacity as learning models with millions of learnable parameters. As a result, they are very prone to overfitting and various convergence difficulties. Consequently, the initial values of a network's weights have paramount importance in the learning process, especially for avoidance of local minima and saddle points.

To address these challenges, prior knowledge in the form of a network's weights can be transferred between models that are later fine-tuned with new data. Azizpour et al. [19] showed that the success of knowledge transfer depends on the similarity between the training dataset of a CNN, and the dataset to which the knowledge is transferred. Given the limited availability of large medical datasets, research on transfer learning in medical imaging is largely focussed on transferring knowledge from general natural images datasets. However, these datasets have very different

properties to medical datasets, including the fact that in medical datasets objects of interest may be very small and boundaries are of paramount importance. Consequently, knowledge transfer between these two domains is not optimal and produces various success rates [19]. In this paper we show that knowledge transfer even between small medical datasets can produce state-of-the-art results with an appropriate network architecture. To the best of our knowledge, this is the first time that deep transfer learning has been applied in the MA detection domain.

The main contributions of this paper are as follows. First, we propose a MA detection method that requires only three stages of analysis. Second, we present a novel CNN with a dedicated architecture for MA detection that does not require hand-crafted features. Third, we show how to successfully transfer knowledge between small datasets in MA domain - an important innovation in this domain as retinal image set characteristics vary between cameras, so that any practically useful method must be capable of simple and reliable retraining.

This paper is organized as follows. The proposed method is described in Section 2. Section 3 describes the datasets and performance metrics used for experiments. In Section 4 the evaluation results are presented and compared with existing approaches. Finally, in Section 5 discussion and conclusions are given.

2. Proposed method

Fig. 1 shows a general overview of the proposed method. It consists of three main stages: preprocessing, patch extraction and pixel-wise classification. The main objective of the preprocessing stage is to remove the non-uniform illumination and redundant data from images. The patch extraction stage prepares data for analysis, whereas the pixel-wise classification is performed by a CNN with a novel architecture.

2.1. Preprocessing

First, we extract the green plane of the fundus image as it provides the highest contrast between foreground structures, such as lesions and vessels, and the background. Since we are only interested in pixels inside a Field-of-View (FOV), we automatically generate a mask for pixels outside the FOV. A mask is generated by applying Otsu thresholding [20] to the green plane of the image. Noisy regions are removed by morphological opening and closing with a structuring element of size five. Next, the image is cropped to the size defined by its FOV to accelerate further processing. Subsequently, the image is resized to the smallest image width of the E-ophtha dataset [18], while maintaining the aspect ratio, using bicubic interpolation. Simultaneously, the same operations are applied to the corresponding annotation image. Finally, each image (I) was preprocessed (I_p) by computing a weighted sum as in Eq. (1):

$$I_p = I \cdot \alpha + I_{Gauss} \cdot \beta + \gamma \quad (1)$$

where $\alpha = 4$ and $\beta = -10$ are weight factors; I_{Gauss} is Gaussian blurred image that was created using filter computed as described in Eq. (2) with $\sigma = 10$; $\gamma = 128$ is a scalar added to each sum.

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \quad (2)$$

All values were determined experimentally. Fig. 3 shows an example preprocessed image.

2.2. Pixel-wise classification

The main goal of this stage is to classify each pixel as either MA or non-MA. We cast pixel-wise classification as a probabilistic classification task, where each pixel can be assigned a continuous

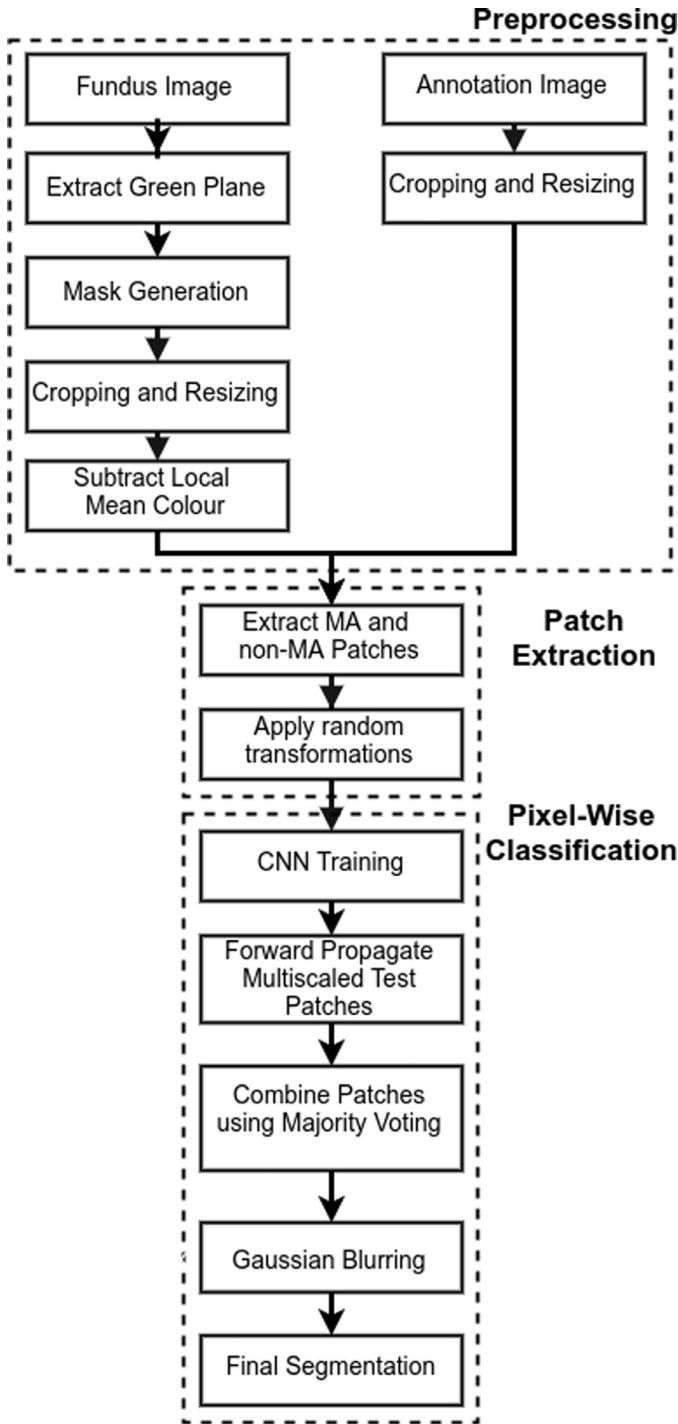


Fig. 1. Overview of the proposed algorithm.

value between 0 (non-MA) and 1 (MA). Compared to other works which perform a binary classification, this learning task is more challenging because the expected output is more complex, hence the underlying data distribution function is harder to model.

The CNN is trained to map an image patch P to the corresponding annotation $A(P)$ for all possible locations within an image. A training sample consists of $S \times S$ sized P and $A(P) : \{P, A(P)\}$.

The goal of training is to learn a mapping $P \rightarrow A(P)$ in the form of a CNN by minimizing

$$L = \sum_{i=1}^N l(A(P)_i, f(P_i; \Theta)) + \Phi(\Theta), \quad (3)$$

where $A(P)_i$ and P_i are the i th annotation patch and i th image patch, N is the number of training samples, $l(\cdot)$ is the loss function, Θ are learning parameters, and $\Phi(\Theta)$ is the regularization term.

2.2.1. Patch generation

At training time, all possible image patches are extracted from each training image using a sliding window approach with 2×2 stride. The patches are divided into two groups: MA patches containing at least 1 MA pixel and non-MA patches consisting of all remaining patches. Both MA and non-MA patches are randomly sampled from the set of all possible patches. Patches that are completely outside the FOV are discarded. Each training sample is subject to random artificial transformations (AT) including rotation, horizontal and vertical reflections with 0.5 probability. The ATs are performed to increase variety in the training set and combat overfitting; they are performed during CNN training so their computational footprint is limited. The proposed method works on a pixel level hence even MA patches consist of more non-MA pixels than MA pixels. As such, MA patches provide both positive and negative training samples. Nevertheless, we added a small set of non-MA patches to the training set to provide network with examples of as many as possible retinal structures (e.g. fovea, optic nerve head) and backgrounds. As a result, the training set consists in 80% of MA patches and in 20% of non-MA patches.

At testing time, all possible image patches from inside of a FOV are extracted. To reconstruct the final image segmentation a voting mechanism is used. Each $A(P)$ produced by the model provides a single vote for all pixels it contains. Given that patches are centred at all possible locations and the $A(P)$ size is $S \times S$, each pixel receives S^2 votes, and a pixel receiving v votes as an MA is assigned a probability of v/S^2 . As a result, a confidence map for pixel MA membership is created.

2.2.2. CNN architecture and training

Inspired by Ronneberger et al. [21], we adopted a fully convolutional approach when designing the CNN. The architecture of the CNN is similar to a convolutional autoencoder: it consists of “contracting” and “expanding” paths. The “contracting” path is used to extract most discriminative features from input (encode the input), whereas the “expanding” path is tasked with recreating and classifying the input by using upscaling and 1×1 convolution operations. Skip connections between the two paths allow for a direct flow of feature maps from earlier to latter layers, which is beneficial for the learning process [22]. Ronneberger et al. [21] designed their fully convolutional neural networks for segmentation of whole images in one pass. As MAs are local features, it is more appropriate here to use a network with a small receptive field and a sliding window approach to processing. Compared with Ronneberger et al. [21], the proposed architecture works on small image patches, incorporates batch normalization (BN) layers and uses different loss function. As MAs occupy a very small proportion of fundus images that feature them, there is a significant class imbalance in the problem domain. To address this we incorporated a Dice coefficient function [23] as a loss function as it effectively handles the overwhelming number of true negatives. The Dice coefficient loss function was used before with CNNs [22] but not in context of MA detection. The training algorithm maximises the Dice loss function which measures the overlap between ground truths y and predicted segmentation \hat{y} . Its values range between 0 (no overlap) and 1 (perfect agreement) and is calculated as

$$DICE = \frac{2 * |y \cap \hat{y}| + \delta}{|y| + |\hat{y}| + \delta} \quad (4)$$

where δ is a small smoothing factor that counteracts against zero value and zero denominator.

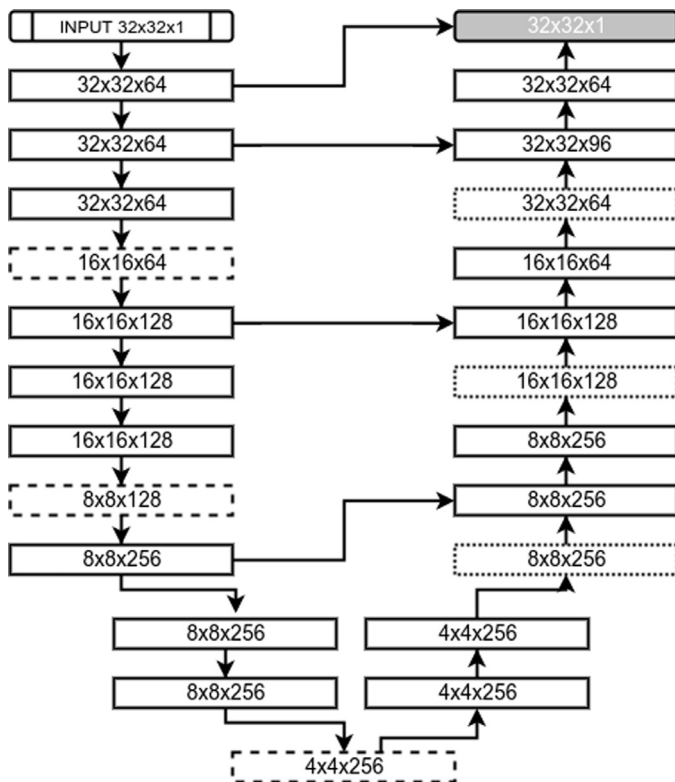


Fig. 2. CNN Architecture. Each block provides the shape of its output. Solid line blocks consists of a convolutional and batch normalization layers. Dashed line blocks correspond to pooling layers. Dotted line blocks represent upsampling layers. The final grey block is the final convolutional layer.

The MA detection domain suffers from a common problem in medical imaging that stems from data scarcity, known as *Covariate Shift*: the distribution of features is different for subsets of training and test datasets which violates the i.i.d. (independent and identically distributed) assumption of many machine learning (ML) algorithms [24]. This may result from the use of different retinal camera systems and/or camera settings. The Covariate Shift in small datasets renders the modelling of true data distribution using ML models virtually impossible. To mitigate this difficulty and make data comparable across features, a normalization technique (shifting data to zero mean and unit variance) is used as a preprocessing step [24]. The same phenomenon occurs during training deep CNNs which are hierarchical in nature and is called *Internal Covariate Shift* [25]. A small change in lower layers can cause a landslide effect in upper layers due to changes in the distribution of upper layer inputs. Ioffe and Szegedy [25] proposed a batch normalization layer that partially alleviates the Internal Covariate Shift by normalizing/whitening data flowing between layers. The use of BN layers in CNNs results in faster convergence (higher learning rates) and better regularization (by constraining layer's inputs, its weights are also indirectly constrained).

The CNN architecture was determined experimentally and is depicted in Fig. 2. It consists of 18 convolutional layers, each followed by a BN layer apart from the final classification layer; three 2×2 max-pooling layers in the “contractive” path and corresponding three 2×2 simple upsampling layers that replicate rows and columns of data in the “expanding” path; 4 skip connections between both paths. Double inputs in the “expanding” path are merged by concatenation. All convolutional layers use 3×3 filters and ReLU activation function [26] apart from the final layer which uses a sigmoid activation function. Weights are updated using stochastic gradient descent with batch size 128 and Adam opti-




Fig. 3. Example image from E-Ophtha dataset. From left to right: original image; preprocessed image.

mization technique [27] with 0.0001 initial learning rate. All training pairs are shuffled between each epoch.

Fine-tuning is a process of training a neural network from a set of pre-defined weights [28]. A traditional approach to fine-tune deep neural networks (DNN) is to train only the final layers of a network using a small learning rate. Similarly to Tajbakhsh et al. [28], it was observed that such approach can provide sub-optimal performance. To find the best ratio between trained and frozen layers, an iterative approach with varying train/freeze ratio was employed on a small dataset.

3. Materials and evaluation

The proposed algorithm was evaluated using most widely used performance metrics and publicly available datasets which are described below.

3.1. Datasets

E-Ophtha dataset [18] consists of 381 compressed images of which 148 have MAs presents and 233 depict healthy FPs. Images were acquired at more than 30 screening centres around France at various resolutions at 45° FOV. There are no separate testing and training datasets provided. The variety of image quality, resolution and illumination conditions makes it the most challenging publicly available dataset. To the best of the authors' knowledge, this is the only public dataset that provides pixel-wise ground truths of MAs.

ROC dataset [29] is composed of 50 training and 50 test compressed images. Images were captured by three different fundus cameras at various resolutions ranging from 768×576 to 1389×1383 at 45° FOV. All images were annotated by four experienced graders. Since test ground truths were never made public and the ROC competition website is inactive [29], only training ground truths are available. 37 images of the training set have at least one MA present, and remaining 13 images present healthy FPs.

DIARETDB1 dataset [30] comprises of 28 training and 61 test uncompressed images acquired at 50° FOV. Each 1500×1152 image was manually annotated for presence of MAs and HEs by four medical experts. The final ground truths were created by fusing all annotations with 75% confidence. 38 FPs have no MAs present whereas remaining 51 FPs have at least one MA.

Since the E-Ophtha dataset does not provide separate train and test sets, it is randomly divided into two sets containing 190 and 191 images respectively. During experimentation 2-fold cross-validation is performed, with each subset alternatively treated as the training or testing set. A similar approach is used with the ROC training dataset, which is split into two sets of 25 images each. DIARETDB1 is explicitly divided into training and testing datasets and we utilise the standard split during experiments. ROC and DIARETDB1 datasets do not provide pixel-wise ground truths however they offer central points and radii of all MAs. Following com-

Table 1
Training data.

Dataset	Nr of training images	Nr training patches
ROC	50	72 481
DIARETDB1	28	40 549
E-Ophtha	381	552 451

mon practice, we use this information to calculate evaluation metrics. All datasets have been acquired using similar FOV (either 45° or 50°). As a result, the downsampling process produces lesions with a common scale. It is important to note that when dealing with images acquired using very different FOVs, the downsampling alone is not enough to successfully normalize lesions and other techniques are necessary (e.g. FOV cropping).

3.2. Evaluation metrics

The free-response ROC (FROC) curve is the most commonly used metric for abnormality detection in medical imaging. It plots per-lesion sensitivity against the average number of false positives per image for different threshold values. In contrast to ROC or specificity-based measures, FROC provides meaningful statistics despite the class imbalance between non-MA and MA pixels in an image. Following common practice we calculate a sensitivity score at seven average false positives per image (FPI) points: 1/8, 1/4, 1/2, 1, 2, 4, 8 [29]. Following common practice, we define lesion as a true positive if at least one pixel overlaps with a corresponding ground truth lesion [12]. We performed Wilcoxon signed ranked tests to estimate the statistical significance of results. Tests were conducted using 255 sensitivity values corresponding to all possible greyscale threshold values produced by tested methods.

4. Experimental results

To assess the performance of the proposed method we performed two sets of experiments. In the first set of experiments we evaluate and compare fine-tuning schemes. In the second, we compare the performance of proposed MA detection technique with other state-of-the-art methods.

The implementation was based on Keras deep learning framework [31] and Tensorflow numerical computation library [32]. The experiments were conducted using a PC with Intel Core i7-6700K CPU, two NVIDIA TitanX graphics cards, and 64GB of RAM.

4.1. Model description

Table 1 shows the amount of training images and patches used for experiments. 10% of the training samples are held back as a validation set and an early stopping criteria is used: training stops when validation error does not improve for 20 epochs. If the validation error does not improve for 10 epochs, the learning rate is reduced by a factor of 0.3. During testing all possible patches are extracted from the FOV and forward propagated through the network. All experiments apart from the E-Ophtha evaluation use a network trained on 354 randomly selected E-Ophtha images, and evaluated on remaining 27 images, as the base model. All parameters were determined empirically based on authors experience or successful deep learning works [15,16,21]. We observe that the proposed approach is robust to changes in parameters' values. The modification of parameters barely affects the final results, however it has a moderate impact on speed of error convergence. We conclude that the system is not sensitive to small parameters change, however such changes can affect the amount of time needed for training.

Table 2
Comparison of fine-tuning schemes.

Fine-tuning scheme	% trainable parameters	Test Dice	FROC score
No fine-tuning	100	0.0376	0
Full fine-tuning	100	0.0271	0.139
Freeze 3	98.44	0.0616	0.195
Freeze 5	94.12	0.0715	0.215
Freeze 5+BN	94.10	0.0257	0.152
Freeze 8	73.96	0.0970	0.218
Freeze 8+BN	73.88	0.0255	0.154
Freeze 11	39.40	0.1030	0.233
Freeze 14	4.85	0.1060	0
Freeze 16	1.24	0.0981	0.109

4.2. Fine-tuning

To find the optimal fine-tuning scheme we performed 10 experiments using ROC training dataset; we randomly divided this into a 25 image training set and 25 image test set, using the same split for all experiments. The base model for fine-tuning was trained on the E-Ophtha dataset as described above. Unless stated otherwise, during fine-tuning the same early stopping and training hyper-parameters were used as in the case of base model training.

Table 2 shows a comparison of all fine-tuning schemes. The Dice metric was calculated on per-pixel basis for the test dataset. In our experiments we applied both “shallow” and “deep” fine-tuning by iteratively freezing more initial layers as proposed by Tajbakhsh et al. [28]. As expected, networks trained from scratch (no fine-tuning) and fully retrained (full fine-tuning) provided the worst results. The network without any fine-tuning did not produce a FROC score because the lowest achieved FPI was just below 0.5, and to calculate the FROC score all seven FROC values are required. For comparison purposes we assign a 0 value to all methods that fail to produce the FROC score. These approaches do not take full advantage of already provided knowledge in the form of a base model. Freezing BN layers results in worse performance compared with the same models when BN layers are trainable. The network with 14 initial layers frozen achieved a comparably high test DICE, which means that it still produced competitive results for all possible pixels. However, the per-lesion evaluation showed that the lowest FPI it managed to reach was around 0.25 which is not enough to calculate a FROC score. As expected, freezing the final most task-specific layers results in decreased performance. We observe that by increasing the number of frozen initial layers, our model accomplishes the best performance by freezing 11 initial layers and training 7 final layers. As a result, all following experiments will use this fine-tuning scheme when transferring knowledge between datasets.

4.3. Microaneurysm detection

Table 3 presents a performance comparison between the proposed method and state-of-the-art methods using the ROC training dataset. The Freeze All method corresponds to a FCNN without any fine-tuning. Compared to other techniques, the proposed algorithm achieves the highest average FROC score of 0.355. Most importantly, it provides much better performance for low FPIs. For comparison purposes, we present the sensitivities at seven high FPIs. Nonetheless, similarly to Niemeijer et al. [29] we think that sensitivity values at FPI higher than 1.08 are of little clinical importance. Consequently, we provide the performance metrics for much lower FPI in Table 4.

Table 5 shows a comparison of MA detection methods using the DIARETDB1 dataset. Consistently with ROC results, the proposed algorithm produces the highest average score of 0.392. Furthermore, the sensitivities for all FPIs are higher than provided

Table 3

The sensitivities at various FPIs using ROC training dataset.

Method	1	2	4	8	12	16	20	Score
Zhou et al. [33]	0.135	0.155	0.232	0.288	0.325	0.370	0.420	0.275 ± 0.099
Javidi et al. [8]	0.130	0.147	0.209	0.287	0.319	0.353	0.383	0.261 ± 0.093
Zhang et al. [11]	0.127	0.150	0.197	0.289	0.31	0.316	0.330	0.246 ± 0.079
Niemeijer et al. [29]	0.072	0.087	0.101	0.121	0.130	0.185	0.210	0.129 ± 0.047
Freeze All	0.090	0.108	0.128	0.139	0.156	0.163	0.177	0.137 ± 0.029
Proposed Method	0.174	0.243	0.306	0.385	0.431	0.461	0.485	0.355 ± 0.109

Table 4

The sensitivities at low FPIs using ROC training dataset.

Method	1/8	1/4	1/2	1	2	4	8	Score
Freeze All	0.028	0.040	0.063	0.090	0.108	0.128	0.139	0.085 ± 0.040
Proposed Method	0.039	0.067	0.141	0.174	0.243	0.306	0.385	0.193 ± 0.116

Table 5


The sensitivities at various FPIs using DIARETDB1 dataset.

Method	1/8	1/4	1/2	1	2	4	8	Score
Seoud et al. [35]	0.139	0.176	0.242	0.318	0.427	0.543	0.639	0.355 ± 0.175
Antal et al. [36]	0.015	0.020	0.027	0.068	0.118	0.204	0.316	0.110 ± 0.105
Adal et al. [37]	0.029	0.036	0.044	0.102	0.205	0.303	0.568	0.184 ± 0.183
Freeze All	0	0	0.039	0.082	0.125	0.163	0.189	0.085 ± 0.071
Proposed Method	0.187	0.246	0.288	0.365	0.449	0.570	0.641	0.392 ± 0.157

Table 6

The sensitivities at various FPIs using E-optha dataset.

Method	1/8	1/4	1/2	1	2	4	8	Score
Veiga et al. [12]	0.110	0.152	0.222	0.307	0.383	0.494	0.629	0.328 ± 0.174
Proposed Method	0.185	0.313	0.465	0.604	0.716	0.801	0.849	0.562 ± 0.233

**Fig. 4.** FROC curves produced by the proposed method. (a) E-Ophtha; (b) DIARETDB1; (c) ROC Training.**Table 7**Wilcoxon signed rank test results. Since $p \ll 0.05$, results are statistically significant.

Compared Methods	p -value
ROC: Proposed method vs Freeze All	1.97×10^{-43}
DIARETDB1: Proposed method vs Freeze All	2.02×10^{-43}

by other methods. To transfer knowledge from the base model to models used with ROC and DIARETDB1 datasets, 11 initial layers of the base model were frozen with remaining 7 trained with new data. Table 6 presents the performance comparison using E-Ophtha dataset. This dataset is much bigger than the previous datasets which results in bigger training datasets. The DNNs benefit from bigger datasets [34] hence the results are better than compared with other datasets. Fig. 4 presents FROC curves produced by the proposed algorithm for all three datasets.

Table 7 shows results of Wilcoxon signed rank tests between the proposed method and Freeze All method for ROC and DIARETDB1 datasets. The null hypothesis is that the proposed method provides similar results to Freeze All method, whereas the alternative hypothesis is that the proposed method provides better results than Freeze All method. In our case, the null and alternative hypotheses can be defined as $H_0: M_P = M_F$ and $H_1: M_P > M_F$, where M_P and M_F are medians of sensitivity values produced by the proposed method and Freeze All method respectively. Following common practice, we set the significance level at 0.05. Wilcoxon signed rank tests show statistically significant improvement in the sensitivity values when using the proposed approach ($p \ll 0.05$).

Fig. 6 presents examples of lesion detection results. The detection results were calculated at 1.08 FPI rate which is regarded as clinically acceptable [29]. We observe that many false positive detections are difficult to discern even for a human eye. Similarly to Kauppi et al. [30] we observe high inter-observer variability be-




Fig. 5. Detection results in presence of common challenges using image regions extracted from E-Ophtha and DIARETDB1. True positives are green circled and false positives are yellow circled. (a) Correct detection of an MA close to a vessel; (b) Correct detection of a subtle MA close to the end of a small vessel; (c) False detection of a small pigmentation change; (d) Dust artefact close to the optic nerve head which is correctly ignored. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

tween human graders, which negatively affects the quality of provided ground truths and trained models.

Fig. 5 shows examples of various challenging detections. Many detection algorithms have to extract and remove vessels first to correctly detect MAs close to vessels. Fig. 5(a) shows that the proposed method can successfully detect MAs very close to vessels. In Fig. 5(b) the MA is almost at the end of a small vessel. Fig. 5(c) presents a false positive example, which is a subtle pigmentation change. DIARETDB1 dataset contains dust artefacts located in exactly the same location across many images. Fig. 5(d) shows that the proposed method correctly ignores such artefact.

5. Discussion

The proposed algorithm achieves better results than state-of-the-art methods in terms of the FROC metric. Most importantly, it provides highest performance at low FPIs which are particularly significant for screening application. An MA detection system for screening purposes does not have to find all MAs, but enough MAs to help a clinician decide if a patient needs referral. As such, we think that the proposed algorithm would prove useful as a component of a DR screening process.

The total time required to process a single image is around 220 s. The majority of this time is spent on forward propagating the large amount of patches through the network. However, during this study we did not concentrate on algorithm's efficiency, hence the implementation is experimental and can be improved. The processing time per image could be drastically reduced if the forward propagation step would be parallelized across multiple devices. This will reduce the per-image processing time by a factor close to the number of used devices.

6. Conclusions

This paper presents a novel MA detection method evaluated using three publicly available datasets. The proposed algorithm uses a novel FCNN architecture with BN layers and Dice coefficient loss function to segment and detect MAs. Compared to other techniques that typically require five computational stages, the proposed method requires only three. Furthermore, we show how to successfully and efficiently transfer knowledge between small datasets in the MA detection domain.

Almost all current MA detection methods rely on human-crafted features, hence their usability and robustness is dependent on the designer's knowledge, experience, and skills. Such systems have to be manually recalibrated due to ever-changing image modalities. The proposed method extracts the most discriminative features for MA detection automatically and proves to be robust against changes in image illumination or contrast. In the future,




Fig. 6. Examples of lesion detection results for E-Ophtha dataset. The probability threshold is set to 0.68 which corresponds to 61.86% per-lesion sensitivity and 1.08 average FPI rate. True positives are green circled, false positives are yellow circled and false negatives are red circled. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

we are planning to parallelize the inference step and reduce the processing time to the range of seconds.

Conflict of interest statement

No potential conflict of interest was reported by the authors.

Acknowledgments

This research was made possible by a Marie Curie grant from the [European Commission](#) in the framework of the REVAMMAD ITN (Initial Training Research network), Project number 316990.

Supplementary material

Supplementary material associated with this article can be found, in the online version, at [10.1016/j.cmpb.2018.02.016](https://doi.org/10.1016/j.cmpb.2018.02.016).

References

- [1] IdF diabetes atlas, 7th edn., Int. Diabetes Fed. (2015).
- [2] N. Cheung, P. Mitchell, T. Wong, Diabetic retinopathy, *Lancet* 376 (9735) (2010) 124–136.
- [3] J. Ding, T.Y. Wong, Current epidemiology of diabetic retinopathy and diabetic macular edema, *Curr. Diabetes Rep.* 12 (4) (2012) 346–354.
- [4] J.W. Yau, S.L. Rogers, R. Kawasaki, E.L. Lamoureux, J.W. Kowalski, T. Bek, S.-J. Chen, J.M. Dekker, A. Fletcher, J. Grauslund, et al., Global prevalence and major risk factors of diabetic retinopathy, *Diabetes Care* 35 (3) (2012) 556–564.
- [5] R.J. Tapp, J.E. Shaw, C.A. Harper, M.P. De Courten, B. Balkau, D.J. McCarty, H.R. Taylor, T.A. Welborn, P.Z. Zimmet, The prevalence of and factors associated with diabetic retinopathy in the Australian population, *Diabetes Care* 26 (6) (2003) 1731–1737.
- [6] L. Guariguata, D. Whiting, I. Hambleton, J. Beagley, U. Linnenkamp, J. Shaw, Global estimates of diabetes prevalence for 2013 and projections for 2035, *Diabetes Res. Clin. Pract.* 103 (2) (2014) 137–149.
- [7] C. Baudoin, B. Lay, J. Klein, Automatic detection of microaneurysms in diabetic fluorescein angiography, *Revue d'épidémiologie et de santé publique* 32 (3–4) (1983) 254–261.
- [8] M. Javidi, H.-R. Pourreza, A. Harati, Vessel segmentation and microaneurysm detection using discriminative dictionary learning and sparse representation, *Comput. Methods Programs Biomed.* 139 (2017) 93–108.
- [9] L.A. Yannuzzi, K.T. Rohrer, L.J. Tindel, R.S. Sobel, M.A. Costanza, W. Shields, E. Zang, Fluorescein angiography complication survey, *Ophthalmology* 93 (5) (1986) 611–617.
- [10] T. Walter, P. Massin, A. Erginay, R. Ordonez, C. Jeulin, J.-C. Klein, Automatic detection of microaneurysms in color fundus images, *Med. Image Anal.* 11 (6) (2007) 555–566.
- [11] B. Zhang, X. Wu, J. You, Q. Li, F. Karray, Detection of microaneurysms using multi-scale correlation coefficients, *Pattern Recognit.* 43 (6) (2010) 2237–2248.
- [12] D. Veiga, N. Martins, M. Ferreira, J. Monteiro, Automatic microaneurysm detection using laws texture masks and support vector machines, *Comput. Methods Biomed. Eng.* (2017) 1–12.
- [13] M. Haloi, Improved microaneurysm detection using deep neural networks, *arXiv:1505.04424* (2015).
- [14] R. Srivastava, L. Duan, D.W. Wong, J. Liu, T.Y. Wong, Detecting retinal microaneurysms and hemorrhages with robustness to the presence of blood vessels, *Comput. Methods Programs Biomed.* 138 (2017) 83–91.
- [15] K. Simonyan, A. Zisserman, Very deep convolutional networks for large-scale image recognition, *arXiv:1409.1556* (2014).
- [16] J. Long, E. Shelhamer, T. Darrell, Fully convolutional networks for semantic segmentation, in: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2015, pp. 3431–3440.
- [17] A. Krizhevsky, I. Sutskever, G.E. Hinton, Imagenet classification with deep convolutional neural networks, in: *Advances in neural information processing systems*, 2012, pp. 1097–1105.
- [18] E. Decencière, G. Cazuguel, X. Zhang, G. Thibault, J.-C. Klein, F. Meyer, B. Marcote, G. Quéllec, M. Lamard, R. Danno, et al., Teleophtha: machine learning and image processing methods for teleophthalmology, *IRBM* 34 (2) (2013) 196–203.
- [19] H. Azizpour, A. Sharif Razavian, J. Sullivan, A. Maki, S. Carlsson, From generic to specific deep representations for visual recognition, in: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops*, 2015, pp. 36–45.
- [20] N. Otsu, A threshold selection method from gray-level histograms, *IEEE Trans. Syst. Man Cybern.* 9 (1) (1979) 62–66.
- [21] O. Ronneberger, P. Fischer, T. Brox, U-net: convolutional networks for biomedical image segmentation, in: *International Conference on Medical Image Computing and Computer-Assisted Intervention*, Springer, 2015, pp. 234–241.
- [22] M. Drozdal, E. Vorontsov, G. Chartrand, S. Kadoury, C. Pal, The importance of skip connections in biomedical image segmentation, in: *International Workshop on Large-Scale Annotation of Biomedical Data and Expert Label Synthesis*, Springer, 2016, pp. 179–187.
- [23] L.R. Dice, Measures of the amount of ecologic association between species, *Ecology* 26 (3) (1945) 297–302.
- [24] J. Quionero-Candela, M. Sugiyama, A. Schwaighofer, N.D. Lawrence, *Dataset Shift in Machine Learning*, The MIT Press, 2009.
- [25] S. Ioffe, C. Szegedy, Batch normalization: accelerating deep network training by reducing internal covariate shift, *arXiv:1502.03167* (2015).
- [26] G.E. Dahl, T.N. Sainath, G.E. Hinton, Improving deep neural networks for lvcsr using rectified linear units and dropout, in: *Acoustics, Speech and Signal Processing (ICASSP)*, 2013 IEEE International Conference on, IEEE, 2013, pp. 8609–8613.
- [27] D. Kingma, J. Ba, Adam: a method for stochastic optimization, (2014).
- [28] N. Tajbakhsh, J.Y. Shin, S.R. Gurudu, R.T. Hurst, C.B. Kendall, M.B. Gotway, J. Liang, Convolutional neural networks for medical image analysis: full training or fine tuning? *IEEE Trans. Med. Imaging* 35 (5) (2016) 1299–1312.
- [29] M. Niemeijer, B. Van Ginneken, M.J. Cree, A. Mizutani, G. Quéllec, C.I. Sánchez, B. Zhang, R. Hornero, M. Lamard, C. Muramatsu, et al., Retinopathy online challenge: automatic detection of microaneurysms in digital color fundus photographs, *IEEE Trans. Med. Imaging* 29 (1) (2010) 185–195.
- [30] T. Kauppi, V. Kalesnykiene, J.-K. Kamarainen, L. Lensu, I. Sorri, A. Raninen, R. Voutilainen, H. Uusitalo, H. Kälviäinen, J. Pietilä, The diaretdb1 diabetic retinopathy database and evaluation protocol, in: *BMVC*, 2007, pp. 1–10.
- [31] F. Chollet, et al., Keras, 2015, (<https://github.com/fchollet/keras>).
- [32] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G.S. Corrado, A. Davis, J. Dean, M. Devin, et al., Tensorflow: large-scale machine learning on heterogeneous distributed systems, *arXiv:1603.04467* (2016).
- [33] W. Zhou, C. Wu, D. Chen, Y. Yi, W. Du, Automatic microaneurysm detection using the sparse principal component analysis-based unsupervised classification method, *IEEE Access* 5 (2017) 2563–2572.
- [34] Y. Bengio, et al., Learning deep architectures for ai, in: *Foundations and trends® in Machine Learning*, 2, 2009, pp. 1–127.
- [35] L. Seoud, T. Hurtut, J. Chelbi, F. Cheriet, J.P. Langlois, Red lesion detection using dynamic shape features for diabetic retinopathy screening, *IEEE Trans. Med. Imaging* 35 (4) (2016) 1116–1126.
- [36] B. Antal, A. Hajdu, An ensemble-based system for microaneurysm detection and diabetic retinopathy grading, *IEEE Trans. Biomed. Eng.* 59 (6) (2012) 1720–1726.
- [37] K.M. Adal, D. Sidibé, S. Ali, E. Chaum, T.P. Karnowski, F. Mériaudeau, Automated detection of microaneurysms using scale-adapted blob analysis and semi-supervised learning, *Comput. Methods Programs Biomed.* 114 (1) (2014) 1–10.